

Problem Statement

Business Context

Business communities in the United States are facing high demand for human resources, but one of the constant challenges is identifying and attracting the right talent, which is perhaps the most important element in remaining competitive. Companies in the United States look for hard-working, talented, and qualified individuals both locally as well as abroad.

The Immigration and Nationality Act (INA) of the US permits foreign workers to come to the United States to work on either a temporary or permanent basis. The act also protects US workers against adverse impacts on their wages or working conditions by ensuring US employers' compliance with statutory requirements when they hire foreign workers to fill workforce shortages. The immigration programs are administered by the Office of Foreign Labor Certification (OFLC).

OFLC processes job certification applications for employers seeking to bring foreign workers into the United States and grants certifications in those cases where employers can demonstrate that there are not sufficient US workers available to perform the work at wages that meet or exceed the wage paid for the occupation in the area of intended employment.

Objective

In FY 2016, the OFLC processed 775,979 employer applications for 1,699,957 positions for temporary and permanent labor certifications. This was a nine percent increase in the overall number of processed applications from the previous year. The process of reviewing every case is becoming a tedious task as the number of applicants is increasing every year.

The increasing number of applicants every year calls for a Machine Learning based solution that can help in shortlisting the candidates having higher chances of VISA approval. OFLC has hired the firm EasyVisa for data-driven solutions. You as a data scientist at EasyVisa have to analyze the data provided and, with the help of a classification model:

- Facilitate the process of visa approvals.
- Recommend a suitable profile for the applicants for whom the visa should be certified or denied based on the drivers that significantly influence the case status.

Data Description

The data contains the different attributes of employee and the employer. The detailed data dictionary is given below.

- case_id: ID of each visa application
- continent: Information of continent the employee
- education_of_employee: Information of education of the employee
- has_job_experience: Does the employee has any job experience? Y= Yes; N = No
- requires_job_training: Does the employee require any job training? Y = Yes; N = No
- no_of_employees: Number of employees in the employer's company
- yr_of_estab: Year in which the employer's company was established
- region_of_employment: Information of foreign worker's intended region of employment in the US.
- prevailing_wage: Average wage paid to similarly employed workers in a specific occupation in the area of intended employment. The purpose of the prevailing wage is to ensure that the foreign worker is not underpaid compared to other workers offering the same or similar service in the same area of employment.
- unit_of_wage: Unit of prevailing wage. Values include Hourly, Weekly, Monthly, and Yearly.
- full_time_position: Is the position of work full-time? Y = Full Time Position; N = Part Time Position
- case_status: Flag indicating if the Visa was certified or denied

Installing and Importing the necessary libraries

In [101]:

```
# Installing the libraries with the specified version.  
!pip install numpy==1.25.2 pandas==1.5.3 scikit-learn==1.5.2 matplotlib==3.7.1 seaborn==0.13.1 xgboost==2.0.3 -q  
--user
```

Note: After running the above cell, kindly restart the notebook kernel and run all cells sequentially from the below.

In [102]:

```
# To help with reading and manipulation of data
import numpy as np
import pandas as pd

# To help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# To split the data
from sklearn.model_selection import train_test_split

# To impute missing values
from sklearn.impute import SimpleImputer

# To build a Random forest classifier
from sklearn.ensemble import RandomForestClassifier

# To tune a model
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV

# To get different performance metrics
import sklearn.metrics as metrics
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    recall_score,
    accuracy_score,
    precision_score,
    f1_score,
)
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score

# Libraries different ensemble classifiers
from sklearn.ensemble import (
    BaggingClassifier,
    RandomForestClassifier,
    AdaBoostClassifier,
    GradientBoostingClassifier
)

from xgboost import XGBClassifier
from sklearn.tree import DecisionTreeClassifier

# To undersample and oversample the data
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler

# To suppress warnings
import warnings
warnings.filterwarnings("ignore")

# To suppress warnings
import warnings

warnings.filterwarnings("ignore")
```

Import Dataset

In [103]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In [104]:

```
path = '/content/drive/My Drive/python/EasyVisa.csv'
df_orig = pd.read_csv(path)
df = df_orig.copy()
```

Overview of the Dataset

View the first and last 5 rows of the dataset

In [105]:

```
df.head()
```

Out[105]:

	case_id	continent	education_of_employee	has_job_experience	requires_job_training	no_of_employees	yr_of_estab	region_of_empl
0	EZYV01	Asia	High School	N	N	14513	2007	
1	EZYV02	Asia	Master's	Y	N	2412	2002	N
2	EZYV03	Asia	Bachelor's	N	Y	44444	2008	
3	EZYV04	Asia	Bachelor's	N	N	98	1897	
4	EZYV05	Africa	Master's	Y	N	1082	2005	

In [106]:

```
df.tail()
```

Out[106]:

	case_id	continent	education_of_employee	has_job_experience	requires_job_training	no_of_employees	yr_of_estab	region
25475	EZYV25476	Asia	Bachelor's	Y	Y	2601	2008	
25476	EZYV25477	Asia	High School	Y	N	3274	2006	
25477	EZYV25478	Asia	Master's	Y	N	1121	1910	
25478	EZYV25479	Asia	Master's	Y	Y	1918	1887	
25479	EZYV25480	Asia	Bachelor's	Y	N	3195	1960	

Understand the shape of the dataset

In [107]:

```
df.shape
```

Out[107]:

```
(25480, 12)
```

Check the data types of the columns for the dataset

In [108]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25480 entries, 0 to 25479
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   case_id          25480 non-null   object  
 1   continent        25480 non-null   object  
 2   education_of_employee  25480 non-null   object  
 3   has_job_experience 25480 non-null   object  
 4   requires_job_training 25480 non-null   object  
 5   no_of_employees    25480 non-null   int64  
 6   yr_of_estab       25480 non-null   int64  
 7   region_of_employment 25480 non-null   object  
 8   prevailing_wage   25480 non-null   float64 
 9   unit_of_wage      25480 non-null   object  
 10  full_time_position 25480 non-null   object  
 11  case_status       25480 non-null   object  
dtypes: float64(1), int64(2), object(9)
memory usage: 2.3+ MB
```

Exploratory Data Analysis (EDA)

Let's check the statistical summary of the data

In [109]:

```
summary = df.describe(include="all");
print(summary)

      case_id continent education_of_employee has_job_experience \
count      25480      25480            25480          25480
unique     25480           6              4             2
top       EZYV25480        Asia        Bachelor's            Y
freq         1        16861            10234          14802
mean        NaN        NaN            NaN            NaN
std         NaN        NaN            NaN            NaN
min        NaN        NaN            NaN            NaN
25%        NaN        NaN            NaN            NaN
50%        NaN        NaN            NaN            NaN
75%        NaN        NaN            NaN            NaN
max        NaN        NaN            NaN            NaN

      requires_job_training no_of_employees yr_of_estab \
count      25480      25480.000000  25480.000000
unique        2             NaN            NaN
top          N             NaN            NaN
freq      22525             NaN            NaN
mean        NaN      5667.043210   1979.409929
std         NaN    22877.928848   42.366929
min        NaN      -26.000000  1800.000000
25%        NaN      1022.000000  1976.000000
50%        NaN      2109.000000  1997.000000
75%        NaN      3504.000000  2005.000000
max        NaN    602069.000000  2016.000000

      region_of_employment prevailing_wage unit_of_wage full_time_position \
count      25480      25480.000000      25480          25480
unique        5             NaN            4             2
top       Northeast             NaN            Year            Y
freq       7195             NaN            22962          22773
mean        NaN      74455.814592            NaN            NaN
std         NaN    52815.942327            NaN            NaN
min        NaN      2.136700            NaN            NaN
25%        NaN      34015.480000            NaN            NaN
50%        NaN      70308.210000            NaN            NaN
75%        NaN      107735.512500            NaN            NaN
max        NaN    319210.270000            NaN            NaN

      case_status
count      25480
unique        2
top       Certified
freq      17018
mean        NaN
std         NaN
min        NaN
25%        NaN
50%        NaN
75%        NaN
max        NaN
```

Fixing the negative values in number of employees columns

In [110]:

```
negative_values = df['no_of_employees'] < 0
negative_rows = df[negative_values]
df['no_of_employees'] = df['no_of_employees'].apply(lambda x: max(x, 0))
print(df['no_of_employees'] < 0)
```

```
0      False
1      False
2      False
3      False
4      False
...
25475     False
25476     False
25477     False
25478     False
25479     False
Name: no_of_employees, Length: 25480, dtype: bool
```

Let's check the count of each unique category in each of the categorical variables

In [111]:

```
# Count unique categories for each categorical variable
category_counts = {col: df[col].value_counts() for col in df.select_dtypes(include='object').columns}
```

```
# Display the counts
```

```
for col, counts in category_counts.items():
    print(f"Counts for {col}:\n{counts}\n")
```

Counts for case_id:

```
case_id
EZYV25480      1
EZYV01         1
EZYV02         1
EZYV03         1
EZYV04         1
...
EZYV13         1
EZYV12         1
EZYV11         1
EZYV10         1
EZYV09         1
Name: count, Length: 25480, dtype: int64
```

Counts for continent:

```
continent
Asia            16861
Europe          3732
North America   3292
South America   852
Africa           551
Oceania          192
Name: count, dtype: int64
```

Counts for education_of_employee:

```
education_of_employee
Bachelor's      10234
Master's        9634
High School     3420
Doctorate       2192
Name: count, dtype: int64
```

Counts for has_job_experience:

```
has_job_experience
Y    14802
N    10678
Name: count, dtype: int64
```

Counts for requires_job_training:

```
requires_job_training
N    22525
Y    2955
Name: count, dtype: int64
```

Counts for region_of_employment:

```
region_of_employment
Northeast      7195
South          7017
West           6586
Midwest        4307
Island          375
Name: count, dtype: int64
```

Counts for unit_of_wage:

```
unit_of_wage
Year           22962
Hour           2157
Week           272
Month          89
Name: count, dtype: int64
```

Counts for full_time_position:

```
full_time_position
Y    22773
N    2707
Name: count, dtype: int64
```

Counts for case_status:

```
case_status
Certified      17018
Denied         8462
Name: count, dtype: int64
```

Univariate Analysis

In [112]:

```
def histogram_boxplot(data, feature, figsize=(15, 10), kde=False, bins=None):
    """
    Boxplot and histogram combined

    data: dataframe
    feature: dataframe column
    figsize: size of figure (default (15,10))
    kde: whether to show the density curve (default False)
    bins: number of bins for histogram (default None)
    """

    f2, (ax_box2, ax_hist2) = plt.subplots(
        nrows=2, # Number of rows of the subplot grid= 2
        sharex=True, # x-axis will be shared among all subplots
        gridspec_kw={"height_ratios": (0.25, 0.75)},
        figsize=figsize,
    ) # creating the 2 subplots
    sns.boxplot(
        data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
    ) # boxplot will be created and a triangle will indicate the mean value of the column
    sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins
    ) if bins else sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2
    ) # For histogram
    ax_hist2.axvline(
        data[feature].mean(), color="green", linestyle="--"
    ) # Add mean to the histogram
    ax_hist2.axvline(
        data[feature].median(), color="black", linestyle="-."
    ) # Add median to the histogram
```

In [113]:

```
# function to create labeled barplots

def labeled_barplot(data, feature, perc=False, n=None):
    """
    Barplot with percentage at the top

    data: dataframe
    feature: dataframe column
    perc: whether to display percentages instead of count (default is False)
    n: displays the top n category levels (default is None, i.e., display all levels)
    """

    total = len(data[feature]) # length of the column
    count = data[feature].nunique()
    if n is None:
        plt.figure(figsize=(count + 1, 5))
    else:
        plt.figure(figsize=(n + 1, 5))

    plt.xticks(rotation=90, fontsize=15)
    ax = sns.countplot(
        data=data,
        x=feature,
        palette="Paired",
        order=data[feature].value_counts().index[:n].sort_values(),
    )

    for p in ax.patches:
        if perc == True:
            label = "{:.1f}%".format(
                100 * p.get_height() / total
            ) # percentage of each class of the category
        else:
            label = p.get_height() # count of each level of the category

        x = p.get_x() + p.get_width() / 2 # width of the plot
        y = p.get_height() # height of the plot


        ax.annotate(
            label,
            (x, y),
            ha="center",
            va="center",
            size=12,
            xytext=(0, 5),
            textcoords="offset points",
        ) # annotate the percentage

    plt.show() # show the plot
```

Observations on education of employee

In [114]:

```
# histogram_boxplot(df, 'education_of_employee', figsize=(15, 10), kde=False, bins=None)
labeled_barplot(df, 'education_of_employee', perc=True, n=None)
```




Obsevation: The vast majority of the employees have a bachelors degree (40.16%), followed by Masters (31.81%), the high school and doctorate contribyte to 13.4 and 8.6% respectively.

Observations on region of employment

In [115]:

```
labeled_barplot(df, 'region_of_employment', perc=True, n=None)
```




Observations: The region of employment is even for the most. 28.2% from the north east and 27.5 % from the South. 25.8% from the west

Observations on job experience

In [116]:

```
labeled_barplot(df, 'has_job_experience', perc=True, n=None)
```




Observation: 58.1 % have experience whereas 41.9% do NOT have experience.

Observations on case status

In [117]:

```
labeled_barplot(df, 'case_status', perc=True, n=None)
```



Observation: 66.8% have been approved, the rest were not approved.

Bivariate Analysis

Creating functions that will help us with further analysis.

In [118]:

```
### function to plot distributions wrt target

def distribution_plot_wrt_target(data, predictor, target):
    fig, axs = plt.subplots(2, 2, figsize=(12, 10))
    target_uniq = data[target].unique()

    axs[0, 0].set_title("Distribution of target for target=" + str(target_uniq[0]))
    sns.histplot(
        data=data[data[target] == target_uniq[0]],
        x=predictor,
        kde=True,
        ax=axs[0, 0],
        color="teal",
        stat="density",
    )

    axs[0, 1].set_title("Distribution of target for target=" + str(target_uniq[1]))
    sns.histplot(
        data=data[data[target] == target_uniq[1]],
        x=predictor,
        kde=True,
        ax=axs[0, 1],
        color="orange",
        stat="density",
    )

    axs[1, 0].set_title("Boxplot w.r.t target")
    sns.boxplot(data=data, x=target, y=predictor, ax=axs[1, 0], palette="gist_rainbow")

    axs[1, 1].set_title("Boxplot (without outliers) w.r.t target")
    sns.boxplot(
        data=data,
        x=target,
        y=predictor,
        ax=axs[1, 1],
        showfliers=False,
        palette="gist_rainbow",
    )

    plt.tight_layout()
    plt.show()
```

In [119]:

```
def stacked_barplot(data, predictor, target):
    """
    Print the category counts and plot a stacked bar chart


    data: dataframe
    predictor: independent variable
    target: target variable
    """
    count = data[predictor].nunique()
    sorter = data[target].value_counts().index[-1]
    tab1 = pd.crosstab(data[predictor], data[target], margins=True).sort_values(
        by=sorter, ascending=False
    )
    print(tab1)
    print("-" * 120)
    tab = pd.crosstab(data[predictor], data[target], normalize="index").sort_values(
        by=sorter, ascending=False
    )
    tab.plot(kind="bar", stacked=True, figsize=(count + 5, 5))
    plt.legend(
        loc="lower left", frameon=False,
    )
    plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
    plt.show()
```

Does higher education increase the chances of visa certification for well-paid jobs abroad?

In [120]:

```
stacked_barplot(df, 'education_of_employee', 'case_status')
```

case_status	Certified	Denied	All
education_of_employee			
All	17018	8462	25480
Bachelor's	6367	3867	10234
High School	1164	2256	3420
Master's	7575	2059	9634
Doctorate	1912	280	2192




Observation: Yes, the higher the education better are the changes of acceptance of visa. As one can see above, the doctorate's (indicated by blue) have higher acceptance rates, followed by Masters, then bachelors and then high school.

How does visa status vary across different continents?

In [121]:


```
distribution_plot_wrt_target(df, 'continent', 'case_status')
```



In [122]:

```
stacked_barplot(df, 'continent', 'case_status')
```

case_status	Certified	Denied	All
continent			
All	17018	8462	25480
Asia	11012	5849	16861
North America	2037	1255	3292
Europe	2957	775	3732
South America	493	359	852
Africa	397	154	551
Oceania	122	70	192




Observation: The % of approved applications was highest in Europe, followed by Africe, Asia, Ocenia and the lowest % is from South America.

Does having prior work experience influence the chances of visa certification for career opportunities abroad?

In [123]:

```
stacked_barplot(df, 'has_job_experience', 'case_status')
```

case_status	Certified	Denied	All
has_job_experience			
All	17018	8462	25480
N	5994	4684	10678
Y	11024	3778	14802




Observation: Yes, have prior job experience increases the chances of getting an approval.

Is the prevailing wage consistent across all regions of the US?

In [124]:

```
distribution_plot_wrt_target(df, 'prevailing_wage', 'region_of_employment')
```




Observation: The average prevailing wage in Midwest and Island are higher. So its not the same.

Does visa status vary with changes in the prevailing wage set to protect both local talent and foreign workers?

In [125]:

```
distribution_plot_wrt_target(df, 'prevailing_wage', 'case_status')
```



Observation: The prevailing wage does NOT seem to be a big factor between certified and denied.





Although the average prevailing wage is slightly higher for Certified cases. Its not very far apart to mention of a distinction, visually

Does the unit of prevailing wage (Hourly, Weekly, etc.) have any impact on the likelihood of visa application certification?


In [126]:

```
distribution_plot_wrt_target(df, 'unit_of_wage', 'case_status')
```

```
stacked_barplot(df, 'unit_of_wage', 'case_status')
```



case_status	Certified	Denied	All
unit_of_wage			
All	17018	8462	25480
Year	16047	6915	22962
Hour	747	1410	2157
Week	169	103	272
Month	55	34	89



Observation: yes, the unit of wage seems to a factor in the case status. The amount of denial for yearly wagers is lesser.

Data Pre-processing

Outlier Check

In [127]:

```
# Calculate Q1, Q3, and IQR for prevailing wage
Q1 = df['prevailing_wage'].quantile(0.25)
Q3 = df['prevailing_wage'].quantile(0.75)
IQR = Q3 - Q1

# Define outlier thresholds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Identify outliers
outliers_wage = df[(df['prevailing_wage'] < lower_bound) | (df['prevailing_wage'] > upper_bound)]
print("outlier counts for wage" , outliers_wage.value_counts())
print(outliers_wage)
```

outlier counts for wage case_id continent education_of_employee has_job_experience require_s_job_training no_of_employees yr_of_estab \ region_of_employment prevailing_wage unit_of_wage f ull_time_position case_status

EZYV9898	Asia	High School	Y	N	18
2001	West	272311.21	Year	Y	Denied
1					
EZYV10151	Asia	Bachelor's	Y	N	3525
1970	Island	289878.68	Year	Y	Denied
1					
EZYV10181	Asia	Doctorate	N	N	2705
1996	South	227441.23	Year	N	Certified
1					
EZYV10331	North America	Master's	Y	N	2535
1975	Northeast	245197.29	Year	N	Certified
1					
EZYV1039	Asia	Master's	Y	N	538
2011	Midwest	222628.84	Year	Y	Certified
1					
..					
EZYV1083	Asia	Doctorate	Y	N	6024
1989	Northeast	274019.43	Year	Y	Certified
1					
EZYV10821	Africa	Bachelor's	N	N	2651
2006	West	268713.50	Year	Y	Certified
1					
EZYV10811	Asia	Bachelor's	Y	N	92780
2012	West	250163.10	Year	Y	Certified
1					
EZYV10706	Asia	High School	N	Y	1856
2003	West	228715.49	Year	Y	Denied
1					
EZYV1069	Africa	Master's	N	N	1486
1992	Northeast	233641.72	Year	Y	Denied
1					
Name: count, Length: 427, dtype: int64					
case_id	continent	education_of_employee	has_job_experience	\	
14	EZYV15	Asia	Master's	Y	
34	EZYV35	Asia	Master's	N	
130	EZYV131	South America	High School	N	
216	EZYV217	Asia	Master's	Y	
221	EZYV222	North America	Doctorate	Y	
...	
25191	EZYV25192	Asia	Master's	N	
25195	EZYV25196	North America	Master's	Y	
25468	EZYV25469	Asia	Bachelor's	N	
25469	EZYV25470	North America	Master's	Y	
25476	EZYV25477	Asia	High School	Y	
requires_job_training	no_of_employees	yr_of_estab	\		
14	Y	15756	2006		
34	N	1809	2010		
130	N	2554	2005		
216	N	1515	2001		
221	Y	2518	2010		
...		
25191	N	4983	2005		
25195	N	47	2001		
25468	N	373	2005		
25469	N	2261	1997		

25476

N

3274

2006

```

region_of_employment  prevailing_wage unit_of_wage full_time_position \
14                  South        220081.73      Year          Y
34                  South        225569.73      Year          N
130                 Midwest     247393.01      Year          Y
216                 Midwest     269321.68      Year          N
221                 South        219529.62      Year          Y
...
25191                ...        ...           ...           ...
25195                Midwest    280482.51      Year          Y
25195                South       234308.77      Year          N
25468                Midwest    272715.74      Year          N
25469                Northeast   273772.47      Year          N
25476                Northeast   279174.79      Year          Y

case_status
14      Certified
34      Certified
130     Certified
216     Certified
221     Certified
...
25191     Denied
25195     Certified
25468     Certified
25469     Certified
25476     Certified

```

[427 rows x 12 columns]

Data Preparation for modeling

In [128]:

```

# Lets prepare the data.
df["case_status"] = df["case_status"].apply(lambda x: 1 if x == "Certified" else 0)
X = df.drop(["case_id"], inplace=True, axis=1)
X = df.drop(["case_status"], axis=1)

y = df["case_status"]

# Apply one-hot encoding to categorical features
categorical_cols = X.select_dtypes(include='object').columns
X = pd.get_dummies(X, columns=categorical_cols, drop_first=True)
print(X.shape)

```

(25480, 21)

In [129]:

```

# Splitting data into training, validation and test set:

# first we split data into 2 parts, say temporary and test
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)

# then we split the temporary set into train and validation
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp)

print(X_train.shape, X_val.shape, X_test.shape)

```

(17836, 21) (3822, 21) (3822, 21)

Model Building

Model Evaluation Criterion

- Choose the primary metric to evaluate the model on
- Elaborate on the rationale behind choosing the metric

First, let's create functions to calculate different metrics and confusion matrix so that we don't have to use the same code repeatedly for each model.

- The `model_performance_classification_sklearn` function will be used to check the model performance of models.
- The `confusion_matrix_sklearn` function will be used to plot the confusion matrix.

In [130]:

```
# defining a function to compute different metrics to check performance of a classification model built using sklearn

def model_performance_classification_sklearn(model, predictors, target):
    """
    Function to compute different metrics to check classification model performance

    model: classifier
    predictors: independent variables
    target: dependent variable
    """

    # predicting using the independent variables
    pred = model.predict(predictors)

    acc = accuracy_score(target, pred) # to compute Accuracy
    recall = recall_score(target, pred) # to compute Recall
    precision = precision_score(target, pred) # to compute Precision
    f1 = f1_score(target, pred) # to compute F1-score

    # creating a dataframe of metrics
    df_perf = pd.DataFrame(
        {"Accuracy": acc, "Recall": recall, "Precision": precision, "F1": f1},
        index=[0],
    )

    return df_perf
```

In [131]:

```
def confusion_matrix_sklearn(model, predictors, target):
    """
    To plot the confusion_matrix with percentages

    model: classifier
    predictors: independent variables
    target: dependent variable
    """

    y_pred = model.predict(predictors)
    cm = confusion_matrix(target, y_pred)
    labels = np.asarray([
        ["{}{:0.0f}{}".format(item) + "\n{}{:0.:2%}{}".format(item / cm.flatten().sum(), item) for item in cm.flatten()]
    ]).reshape(2, 2)

    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=labels, fmt="")
    plt.ylabel("True label")
    plt.xlabel("Predicted label")
```

Defining scorer to be used for cross-validation and hyperparameter tuning

Let us choose the **F1 score** since that is a balance between precision and recall. We don't want people being denied visa (when they truly deserve) and also we don't want them approved when they are not supposed to be. (This might be a security issue in the USA)

In [132]:

```
scorer = metrics.make_scorer(metrics.f1_score)
```

We are now done with pre-processing and evaluation criterion, so let's start building the model.

Model building with Original data

In [133]:

```
models = [] # Empty list to store all the models
results = [] # Empty list to store all model's CV scores
names = [] # Empty list to store name of the models

# Appending models into the list
models.append(("Bagging", BaggingClassifier(random_state=42)))
models.append(("Random forest", RandomForestClassifier(random_state=42)))
models.append(("GBM", GradientBoostingClassifier(random_state=42)))
models.append(("Adaboost", AdaBoostClassifier(random_state=42)))
models.append(("dtree", DecisionTreeClassifier(random_state=42)))
```

In [134]:

```
print(X_train.shape, X_val.shape, X_test.shape)
print(y_train.shape, y_val.shape, y_test.shape)
```

```
(17836, 21) (3822, 21) (3822, 21)
(17836,) (3822,) (3822,)
```

In [135]:

```
# Lets cross validate
print("\n" "Cross-Validation performance on training dataset for all the 5 models:" "\n")

#keeping the n_splits low since the running is taking a long time and also crashing the runtime engine sometimes.
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for name, model in models:
    cv_result = cross_val_score(
        estimator=model, X=X_train, y=y_train, scoring = scorer, cv=kfold
    )
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean()))
    model.fit(X_train,y_train)
    perf = model_performance_classification_sklearn(model, X_val, y_val)
    conf = confusion_matrix_sklearn(model, X_val, y_val)
    print("\n" "Performance: Various Parameters" "\n")
    print(perf)
    print("\n" "Performance: Confusion Matrix" "\n")
    print(conf)
    print("\n")
```

Cross-Validation performance on training dataset for all the 5 models:

Bagging: 0.776902593316487

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.694139	0.779083	0.766769	0.772877

Performance: Confusion Matrix

None

Random forest: 0.8062327459069424

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.72449	0.840971	0.768432	0.803067

Performance: Confusion Matrix

None

GBM: 0.824619206135037

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.744113	0.871524	0.773913	0.819823

Performance: Confusion Matrix

None

Adaboost: 0.8200334908816742

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.731031	0.8868	0.753913	0.814975

Performance: Confusion Matrix

None




dtree: 0.7431976564222731



Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.657509	0.747356	0.741835	0.744585

Performance: Confusion Matrix

None





In [136]:


```
#Plotting boxplots for CV scores of all models defined above
fig = plt.figure(figsize=(10, 7))

fig.suptitle("Algorithm Comparison")
ax = fig.add_subplot(111)

plt.boxplot(results)
ax.set_xticklabels(names)

plt.show()
```

Algorithm Comparison



Model Building with Oversampled data

Lets check for class imbalance.

In [137]:

```
# Checking class balance for whole data, train set, validation set, and test set
```

```
print("Target value ratio in y")
print(y.value_counts(1))
print("*" * 80)
print("Target value ratio in y_train")
print(y_train.value_counts(1))
print("*" * 80)
print("Target value ratio in y_val")
print(y_val.value_counts(1))
print("*" * 80)
print("Target value ratio in y_test")
print(y_test.value_counts(1))
print("*" * 80)
```

```
Target value ratio in y
```

```
case_status
1 0.667896
0 0.332104
```

```
Name: proportion, dtype: float64
```

```
*****
```

```
Target value ratio in y_train
```

```
case_status
1 0.667919
0 0.332081
```

```
Name: proportion, dtype: float64
```

```
*****
```

```
Target value ratio in y_val
```

```
case_status
1 0.667975
0 0.332025
```

```
Name: proportion, dtype: float64
```

```
*****
```

```
Target value ratio in y_test
```

```
case_status
1 0.667713
0 0.332287
```

```
Name: proportion, dtype: float64
```

```
*****
```

In [138]:

```
print("Before OverSampling, counts of label '1': {}".format(sum(y_train == 1)))
print("Before OverSampling, counts of label '0': {} \n".format(sum(y_train == 0)))
```

```
# Synthetic Minority Over Sampling Technique
```

```
sm = SMOTE(sampling_strategy=1, k_neighbors=5, random_state=42)
X_train_over, y_train_over = sm.fit_resample(X_train, y_train)
```

```
print("After OverSampling, counts of label '1': {}".format(sum(y_train_over == 1)))
print("After OverSampling, counts of label '0': {} \n".format(sum(y_train_over == 0)))
```

```
print("After OverSampling, the shape of train_X: {}".format(X_train_over.shape))
print("After OverSampling, the shape of train_y: {} \n".format(y_train_over.shape))
```

```
Before OverSampling, counts of label '1': 11913
Before OverSampling, counts of label '0': 5923
```

```
After OverSampling, counts of label '1': 11913
After OverSampling, counts of label '0': 11913
```

```
After OverSampling, the shape of train_X: (23826, 21)
After OverSampling, the shape of train_y: (23826,)
```

In [139]:

```
models = [] # Empty list to store all the models
results = [] # Empty list to store all model's CV scores
names = [] # Empty list to store name of the models

# Appending models into the list
models.append(("Bagging", BaggingClassifier(random_state=42)))
models.append(("Random forest", RandomForestClassifier(random_state=42)))
models.append(("GBM", GradientBoostingClassifier(random_state=42)))
models.append(("Adaboost", AdaBoostClassifier(random_state=42)))
models.append(("dtree", DecisionTreeClassifier(random_state=42)))

# loop through all models to get the mean cross validated score
print("\n" "Cross-Validation performance on training dataset using oversampled data:" "\n")
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)

for name, model in models:
    cv_result = cross_val_score(
        estimator=model, X=X_train_over, y=y_train_over, scoring = scorer, cv=kfold
    )
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean()))
    model.fit(X_train_over,y_train_over)
    perf = model_performance_classification_sklearn(model, X_val, y_val)
    conf = confusion_matrix_sklearn(model, X_val, y_val)
    print("\n" "Performance: Various Parameters" "\n")
    print(perf)
    print("\n" "Performance: Confusion Matrix" "\n")
    print(conf)
    print("\n")
```

Cross-Validation performance on training dataset using oversampled data:

Bagging: 0.7615542807808515

Performance: Various Parameters

Accuracy	Recall	Precision	F1
0 0.693093	0.762241	0.774682	0.768411

Performance: Confusion Matrix

None

Random forest: 0.7927067506689416

Performance: Various Parameters

Accuracy	Recall	Precision	F1
0 0.719257	0.817861	0.774481	0.79558

Performance: Confusion Matrix

None

GBM: 0.8050749813829021

Performance: Various Parameters

Accuracy	Recall	Precision	F1
0 0.743328	0.854289	0.78172	0.816395

Performance: Confusion Matrix

None

Adaboost: 0.7908741049483738

Performance: Various Parameters

Accuracy	Recall	Precision	F1
0 0.731816	0.855464	0.769014	0.809939

Performance: Confusion Matrix

None




dtree: 0.7301740046260312



Performance: Various Parameters

Accuracy	Recall	Precision	F1
0 0.650183	0.732863	0.740697	0.736759

Performance: Confusion Matrix

None





In [140]:


```
# Plotting boxplots for CV scores of all models defined above
fig = plt.figure(figsize=(10, 7))

fig.suptitle("Algorithm Comparison")
ax = fig.add_subplot(111)

plt.boxplot(results)
ax.set_xticklabels(names)

plt.show()
```

Algorithm Comparison



Model Building with Undersampled data

In [141]:

```
rus = RandomUnderSampler(random_state=42, sampling_strategy=1)
X_train_un, y_train_un = rus.fit_resample(X_train, y_train)

print("Before UnderSampling, counts of label '1': {}".format(sum(y_train == 1)))
print("Before UnderSampling, counts of label '0': {} \n".format(sum(y_train == 0)))

print("After UnderSampling, counts of label '1': {}".format(sum(y_train_un == 1)))
print("After UnderSampling, counts of label '0': {} \n".format(sum(y_train_un == 0)))

print("After UnderSampling, the shape of train_X: {}".format(X_train_un.shape))
print("After UnderSampling, the shape of train_y: {} \n".format(y_train_un.shape))
```

Before UnderSampling, counts of label '1': 11913
Before UnderSampling, counts of label '0': 5923

After UnderSampling, counts of label '1': 5923
After UnderSampling, counts of label '0': 5923

After UnderSampling, the shape of train_X: (11846, 21)
After UnderSampling, the shape of train_y: (11846,)

In [142]:

```
models = [] # Empty list to store all the models
results = [] # Empty list to store all model's CV scores
names = [] # Empty list to store name of the models

# Appending models into the list
models.append(("Bagging", BaggingClassifier(random_state=42)))
models.append(("Random forest", RandomForestClassifier(random_state=42)))
models.append(("GBM", GradientBoostingClassifier(random_state=42)))
models.append(("Adaboost", AdaBoostClassifier(random_state=42)))
models.append(("dtree", DecisionTreeClassifier(random_state=42)))

# loop through all models to get the mean cross validated score
print("\n" "Cross-Validation performance on training dataset using undersampled data:" "\n")
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for name, model in models:
    cv_result = cross_val_score(estimator=model, X=X_train_un, y=y_train_un, scoring = scorer, cv=kfold)
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean()))
    model.fit(X_train_un,y_train_un)
    perf = model_performance_classification_sklearn(model, X_val, y_val)
    conf = confusion_matrix_sklearn(model, X_val, y_val)
    print("\n" "Performance: Various Parameters" "\n")
    print(perf)
    print("\n" "Performance: Confusion Matrix" "\n")
    print(conf)
    print("\n")
```

Cross-Validation performance on training dataset using undersampled data:

Bagging: 0.648227026875054

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.644689	0.609087	0.81201	0.696061

Performance: Confusion Matrix

None

Random forest: 0.6883663664092916

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.68001	0.681943	0.809015	0.740064

Performance: Confusion Matrix

None

GBM: 0.7131226753112041

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.70853	0.732863	0.812419	0.770593

Performance: Confusion Matrix

None

Adaboost: 0.7048541754838573

Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.694924	0.724246	0.800087	0.76028

Performance: Confusion Matrix

None




dtree: 0.6213718649289948



Performance: Various Parameters

	Accuracy	Recall	Precision	F1
0	0.623234	0.636898	0.760168	0.693095

Performance: Confusion Matrix

None





Hyperparameter Tuning

The 3 models that will be used for Hyperparameter tuning

RandomForest - with oversampling

GBM - with oversampling

AdaBoost - with oversampling.

Best practices for hyperparameter tuning in AdaBoost:

n_estimators :

- Start with a specific number (50 is used in general) and increase in steps: 50, 75, 85, 100
- Use fewer estimators (e.g., 50 to 100) if using complex base learners (like deeper decision trees)
- Use more estimators (e.g., 100 to 150) when learning rate is low (e.g., 0.1 or lower)
- Avoid very high values unless performance keeps improving on validation

learning_rate :

- Common values to try: 1.0, 0.5, 0.1, 0.01
- Use 1.0 for faster training, suitable for fewer estimators
- Use 0.1 or 0.01 when using more estimators to improve generalization
- Avoid very small values (< 0.01) unless you plan to use many estimators (e.g., >500) and have sufficient data

In [143]:

```
%%time

# defining model
model = AdaBoostClassifier(random_state=42)

# Parameter grid to pass in RandomSearchCV
param_grid = {
    "n_estimators": [50, 75, 85, 100],
    "learning_rate": [0.01, 0.1, 0.5, 1],
    "estimator": [DecisionTreeClassifier(max_depth=1, random_state=42), DecisionTreeClassifier(max_depth=2, random_state=42), DecisionTreeClassifier(max_depth=3, random_state=42)],
    }
}

#Calling RandomizedSearchCV
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

randomized_cv = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    n_iter=7,
    scoring=scorer,
    cv=kfold,
    random_state=42
)

#Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_over,y_train_over)

# Print the best combination of parameters
adaboost_bestparams = randomized_cv.best_params_
adaboost_bestmodel = randomized_cv.best_estimator_
print("Best parameters: ", adaboost_bestparams)
print("Best model: ", adaboost_bestmodel)

adaboost_train = model_performance_classification_sklearn(adaboost_bestmodel, X_train_over, y_train_over)
adaboost_val = model_performance_classification_sklearn(adaboost_bestmodel, X_val, y_val)

print("\n" "Train Performance:" "\n")
print(adaboost_train)
print("\n" "Validation Performance:" "\n")
print(adaboost_val)
print("\n")
```

```
Best parameters: {'n_estimators': 100, 'learning_rate': 0.5, 'estimator': DecisionTreeClassifier(max_depth=3, random_state=42)}
Best model: AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=3, random_state=42),
                                learning_rate=0.5, n_estimators=100, random_state=42)
```

Train Performance:

	Accuracy	Recall	Precision	F1
0	0.791908	0.856879	0.758339	0.804603

Validation Performance:

	Accuracy	Recall	Precision	F1
0	0.735217	0.852722	0.773907	0.811405

CPU times: user 1min 3s, sys: 78.1 ms, total: 1min 3s
Wall time: 1min 3s

Best practices for hyperparameter tuning in Random Forest:

n_estimators :

- Start with a specific number (50 is used in general) and increase in steps: 50, 75, 100, 125
- Higher values generally improve performance but increase training time
- Use 100-150 for large datasets or when variance is high

min_samples_leaf :

- Try values like: 1, 2, 4, 5, 10
- Higher values reduce model complexity and help prevent overfitting
- Use 1–2 for low-bias models, higher (like 5 or 10) for more regularized models
- Works well in noisy datasets to smooth predictions

max_features :

- Try values: "sqrt" (default for classification), "log2", None, or float values (e.g., 0.3, 0.5)
- "sqrt" balances between diversity and performance for classification tasks
- Lower values (e.g., 0.3) increase tree diversity, reducing overfitting
- Higher values (closer to 1.0) may capture more interactions but risk overfitting

max_samples (for bootstrap sampling):

- Try float values between 0.5 to 1.0 or fixed integers
 - Use 0.6–0.9 to introduce randomness and reduce overfitting
 - Smaller values increase diversity between trees, improving generalization
-

In [144]:

```
%time
# defining model
model = RandomForestClassifier(random_state=42)

# Parameter grid to pass in RandomSearchCV
param_grid = {
    "n_estimators": [50, 75, 100, 125], ## Complete the code to set the number of estimators.
    "min_samples_leaf": [2, 4, 5], ## Complete the code to set the minimum number of samples in the leaf node.
    "max_features": [0.3, 0.5], ## Complete the code to set the maximum number of features.
    "max_samples": [0.6, 0.9], ## Complete the code to set the maximum number of samples.
}

#Calling RandomizedSearchCV
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
randomized_cv = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    n_iter=7,
    scoring=scorer,
    cv=kfold,
    random_state=42
)

#Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_over, y_train_over)
randomForest_bestparams = randomized_cv.best_params_
randomForest_bestmodel = randomized_cv.best_estimator_
print("Best parameters: ", randomForest_bestparams)
print("Best model: ", randomForest_bestmodel)

# perf = model_performance_classification_sklearn(randomForest_bestmodel, X_val, y_val)
# conf = confusion_matrix_sklearn(randomForest_bestmodel, X_val, y_val)
randomForest_trained = model_performance_classification_sklearn(randomForest_bestmodel, X_train_over, y_train_over)
randomForest_val = model_performance_classification_sklearn(randomForest_bestmodel, X_val, y_val)
print("\n" "Performance: " "\n")
print("\n" "Train Performance: " "\n")
print(randomForest_trained)
print("\n" "Validation Performance: " "\n")
print(randomForest_val)
print("\n")
```

Best parameters: {'n_estimators': 50, 'min_samples_leaf': 2, 'max_samples': 0.9, 'max_features': 0.3}
Best model: RandomForestClassifier(max_features=0.3, max_samples=0.9, min_samples_leaf=2, n_estimators=50, random_state=42)

Performance:

Train Performance:

	Accuracy	Recall	Precision	F1
0	0.92781	0.953916	0.906582	0.929647

Validation Performance:

	Accuracy	Recall	Precision	F1
0	0.732601	0.835879	0.779686	0.806805

CPU times: user 1min 28s, sys: 91.7 ms, total: 1min 28s

Wall time: 1min 28s

Best practices for hyperparameter tuning in Gradient Boosting:

n_estimators :

- Start with 100 (default) and increase: 100, 200, 300, 500
- Typically, higher values lead to better performance, but they also increase training time
- Use 200–500 for larger datasets or complex problems
- Monitor validation performance to avoid overfitting, as too many estimators can degrade generalization

learning_rate :

- Common values to try: 0.1, 0.05, 0.01, 0.005
- Use lower values (e.g., 0.01 or 0.005) if you are using many estimators (e.g., > 200)
- Higher learning rates (e.g., 0.1) can be used with fewer estimators for faster convergence
- Always balance the learning rate with n_estimators to prevent overfitting or underfitting

subsample :

- Common values: 0.7, 0.8, 0.9, 1.0
- Use a value between 0.7 and 0.9 for improved generalization by introducing randomness
- 1.0 uses the full dataset for each boosting round, potentially leading to overfitting
- Reducing subsample can help reduce overfitting, especially in smaller datasets

max_features :

- Common values: "sqrt", "log2", or float (e.g., 0.3, 0.5)
 - "sqrt" (default) works well for classification tasks
 - Lower values (e.g., 0.3) help reduce overfitting by limiting the number of features considered at each split
-

In [145]:

```
%%time
# defining model
model = GradientBoostingClassifier(random_state=1)

param_grid={
    "n_estimators": [50, 100, 150, 200, 250],
    "learning_rate": [0.1, 0.05, 0.01, 0.005],
    "subsample": [0.7, 0.8, 0.9, 1.0],
    "max_features": [0.3, 0.5]
}

#Calling RandomizedSearchCV
kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
randomized_cv = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    n_iter=7,
    scoring=scorer,
    cv=kfold,
    random_state=1
)

#Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_over, y_train_over)
gbBoost_bestparams = randomized_cv.best_params_
gbBoost_bestmodel = randomized_cv.best_estimator_
print("Best GD Boost parameters: ", gbBoost_bestparams)
print("Best GD Boost model: ", gbBoost_bestmodel)

gbm_trained = model_performance_classification_sklearn(gbBoost_bestmodel, X_train_over, y_train_over)
gbm_val = model_performance_classification_sklearn(gbBoost_bestmodel, X_val, y_val)
print("\n" "Performance: on trained data" "\n")
print(gbm_trained)
print("\n" "Performance: on validated data" "\n")
print(gbm_val)
print("\n")
```

Best GD Boost parameters: {'subsample': 1.0, 'n_estimators': 250, 'max_features': 0.3, 'learning_rate': 0.1}
Best GD Boost model: GradientBoostingClassifier(max_features=0.3, n_estimators=250, random_state=1)

Performance: on trained data

	Accuracy	Recall	Precision	F1
0	0.802233	0.863007	0.769478	0.813563

Performance: on validated data

	Accuracy	Recall	Precision	F1
0	0.74202	0.857031	0.778925	0.816113

CPU times: user 51.3 s, sys: 69.3 ms, total: 51.3 s
Wall time: 51.4 s

Model Performance Summary and Final Model Selection

We will go with the GBM model since they gave the best F1 scores.

The details have been provided in the table(s) below

In [146]:

```
# training performance comparison
models_train_comp_df = pd.concat([
    gbm_trained.T,
    adaboost_train.T,
    randomForest_trained.T
], axis=1,
)
models_train_comp_df.columns = [
    "GBM (oversampled)",
    "AdaBoost (oversampled)",
    "Random forest (oversampled)"
]
print("Training performance comparison:")
models_train_comp_df
```

Training performance comparison:

Out[146]:

	GBM (oversampled)	AdaBoost (oversampled)	Random forest (oversampled)
Accuracy	0.802233	0.791908	0.927810
Recall	0.863007	0.856879	0.953916
Precision	0.769478	0.758339	0.906582
F1	0.813563	0.804603	0.929647

In [147]:

```
# validation performance comparison
models_val_comp_df = pd.concat([
    gbm_val.T,
    adaboost_val.T,
    randomForest_val.T
], axis=1,
)
models_val_comp_df.columns = [
    "GBM (oversampled)",
    "AdaBoost (oversampled)",
    "Random forest (oversampled)"
]
print("Validation performance comparison:")
models_val_comp_df
```

Validation performance comparison:

Out[147]:

	GBM (oversampled)	AdaBoost (oversampled)	Random forest (oversampled)
Accuracy	0.742020	0.735217	0.732601
Recall	0.857031	0.852722	0.835879
Precision	0.778925	0.773907	0.779686
F1	0.816113	0.811405	0.806805

As we can see that the gradient boosting with oversampling gives consistent results in training and validation.

The F1 scores are 0.813 and 0.816.

So we will go with that.

In [148]:

```
final_model = gbBoost_bestmodel
```


In [149]:

```
production_model = model_performance_classification_sklearn(final_model, X_test, y_test)
print(production_model)
print("\n")

feature_names = X_train.columns
importances = final_model.feature_importances_
indices = np.argsort(importances)

plt.figure(figsize=(12, 12))
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="green", align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```

	Accuracy	Recall	Precision	F1
0	0.73888	0.852273	0.777897	0.813388



Actionable Insights and Recommendations

[Power Ahead](#)

The gradient boosting model gives the best F1 scores.

- The F1 scores are pretty consistent across training, validation and test which shows very clearly that we have not overfit.

In order to increase the likeliness of a visa certification (approval)

1. Have a PhD or a masters degree
2. Being from Europe or Africa increases the chances.
3. Having prior work experience also increases the chances
4. Having yearly wage also increases the chances.

We can conclude that we need to let the visa aspirants know that they need to study hard and get at least a bachelors and have work experience. This would surely increase their chances of getting a visa approved.